

= WinFS =

WinFS ( short for Windows Future Storage ) is the code name for a canceled data storage and management system project based on relational databases , developed by Microsoft and first demonstrated in 2003 as an advanced storage subsystem for the Microsoft Windows operating system , designed for persistence and management of structured , semi @-@ structured as well as unstructured data .

WinFS includes a relational database for storage of information , and allows any type of information to be stored in it , provided there is a well defined schema for the type . Individual data items could then be related together by relationships , which are either inferred by the system based on certain attributes or explicitly stated by the user . As the data has a well defined schema , any application can reuse the data ; and using the relationships , related data can be effectively organized as well as retrieved . Because the system knows the structure and intent of the information , it can be used to make complex queries that enable advanced searching through the data and aggregating various data items by exploiting the relationships between them .

While WinFS and its shared type schema make it possible for an application to recognize the different data types , the application still has to be coded to render the different data types . Consequently , it would not allow development of a single application that can view or edit all data types ; rather what WinFS enables applications to understand is the structure of all data and extract the information that it can use further . When WinFS was introduced at the 2003 Professional Developers Conference , Microsoft also released a video presentation , named IWish , showing mockup interfaces that showed how applications would expose interfaces that take advantage of a unified type system . The concepts shown in the video ranged from applications using the relationships of items to dynamically offer filtering options to applications grouping multiple related data types and rendering them in a unified presentation .

WinFS was billed as one of the pillars of the " Longhorn " wave of technologies , and would ship as part of the next version of Windows . It was subsequently decided that WinFS would ship after the release of Windows Vista , but those plans were shelved in June 2006 , with some of its component technologies being integrated into upcoming releases of ADO.NET and Microsoft SQL Server . While it was then assumed by observers that WinFS was finished as a project , in November 2006 Steve Ballmer announced that WinFS was still in development , though it was not clear how the technology was to be delivered . Several components of the last Integrated Storage Initiative project , Microsoft Semantic Engine , presented at Microsoft PDC 2009 , have been integrated back into the SQL Server " Denali " . At the 2010 SQL Server PASS Community Summit , the forthcoming version of SQL Server ( " Denali " ) was shown , which seems to incorporate many of the WinFS ideas .

= = Motivation = =

Many filesystems found on common operating systems , including the NTFS filesystem which is used in modern versions of Microsoft Windows , store files and other objects only as a stream of bytes , and have little or no information about the data stored in the files . Such file systems also provide only a single way of organizing the files , namely via directories and file names .

Because a file system has no knowledge about the data it stores , applications tend to use their own , often proprietary , file formats . This hampers sharing of data between multiple applications . It becomes difficult to create an application which processes information from multiple file types , because the programmers have to understand the structure and semantics of all the files . Using common file formats is a workaround to this problem but not a universal solution ; there is no guarantee that all applications will use the format . Data with standardized schema , such as XML documents and relational data fare better , as they have a standardized structure and run @-@ time requirements .

Also , a traditional file system can retrieve and search data based only on the filename , because the only knowledge it has about the data is the name of the files that store the data . A better solution is to tag files with attributes that describe them . Attributes are metadata about the files such

as the type of file ( such as document , picture , music , creator , etc . ) . This allows files to be searched for by their attributes , in ways not possible using a folder hierarchy , such as finding " pictures which have person X " . The attributes can be recognizable by either the file system natively , or via some extension . Desktop search applications take this concept a step further . They extract data , including attributes , from files and index it . To extract the data , they use a filter for each file format . This allows for searching based on both the file 's attributes and the data in it .

However , this still does not help in managing related data , as disparate items do not have any relationships defined . For example , it is impossible to search for " the phone numbers of all persons who live in Acapulco and each have more than 100 appearances in my photo collection and with whom I have had e @-@ mail within last month " . Such a search could not be done unless it is based on a data model which has both the semantics as well as relationships of data defined . WinFS aims to provide such a data model and the runtime infrastructure that can be used to store the data , as well as the relationships between data items according to the data model , doing so at a satisfactory level of performance .

= = Overview = =

WinFS natively recognizes different types of data , such as picture , e @-@ mail , document , audio , video , calendar , contact , rather than just leaving them as raw unanalyzed bytestreams ( as most file systems do ) . Data stored and managed by the system are instances of the data type recognized by the WinFS runtime . The data are structured by means of properties . For example , an instance of a résumé type will surface the data by exposing properties , such as Name , Educational Qualification , Experience . Each property may be a simple type ( strings , integers , dates ) or complex types ( contacts ) . Different data types expose different properties . Besides that , WinFS also allows different data instances to be related together ; such as a document and a contact can be related by an Authored By relationship . Relationships are also exposed as properties ; for example if a document is related to a contact by a Created By relationship , then the document will have a Created By property . When it is accessed , the relationship is traversed and the related data returned . By following the relations , all related data can be reached . WinFS promotes sharing of data between applications by making the data types accessible to all applications , along with their schemas . When an application wants to use a WinFS type , it can use the schema to find the data structure and can use the information . So , an application has access to all data on the system even though the developer did not have to write parsers to recognize the different data formats . It can also use relationships and related data to create dynamic filters to present the information the application deals with . The WinFS API further abstracts the task of accessing data . All WinFS types are exposed as .NET objects with the properties of the object directly mapping to the properties of the data type . Also , by letting different applications which deal with the same data share the same WinFS data instance rather than storing the same data in different files , the hassles of synchronizing the different stores when the data change are removed . Thus WinFS can reduce redundancies .

Access to all the data in the system allows complex searches for data across all the data items managed by WinFS . In the example used above ( " the phone numbers of all persons who live in Acapulco and each have more than 100 appearances in my photo collection and with whom I have had e @-@ mail within last month " ) , WinFS can traverse the subject relationship of all the photos to find the contact items . Similarly , it can filter all emails in last month and access the communicated with relation to reach the contacts . The common contacts can then be figured out from the two sets of results and their phone number retrieved by accessing the suitable property of the contact items .

In addition to fully schematized data ( like XML and relational data ) , WinFS supports semi @-@ structured data ( such as images , which have an unstructured bitstream plus structured metadata ) as well as unstructured data ( such as files ) as well . It stores the unstructured components as files while storing the structured metadata in the structured store . Internally , WinFS uses a relational database to manage data . It does not limit the data to belonging to any particular data model . The

WinFS runtime maps the schema to a relational modality , by defining the tables it will store the types in and the primary keys and foreign keys that would be required to represent the relationships . WinFS includes mappings for object and XML schemas by default . Mappings for other schemas must be specified . Object schemas are specified in XML ; WinFS generates code to surface the schemas as .NET classes . ADO.NET can be used to directly specify the relational schema , though a mapping to the object schema must be provided to surface it as classes . Relationship traversals are performed as joins on these tables . WinFS also automatically creates indexes on these tables , to enable fast access to the information . Indexing speeds up joins significantly , and traversing relationships to retrieve related data is performed very fast . Indexes are also used during information search ; searching and querying use the indexes to quickly complete the operations , much like desktop search systems .

= = Development = =

The development of WinFS is an extension to a feature which was initially planned in the early 1990s . Dubbed Object File System , it was supposed to be included as part of Cairo . OFS was supposed to have powerful data aggregation features , but the Cairo project was shelved , and with it OFS . However , later during the development of COM , a storage system , called Storage + , based on then @-@ upcoming SQL Server 8 @.@ 0 , was planned , which was slated to offer similar aggregation features . This , too , never materialized , and a similar technology , Relational File System ( RFS ) , was conceived to be launched with SQL Server 2000 . However , SQL Server 2000 ended up being a minor upgrade to SQL Server 7 @.@ 0 and RFS was not implemented .

But the concept was not scrapped . It just morphed into WinFS . WinFS was initially planned for inclusion in Windows Vista , and build 4051 of Windows Vista , then called by its codename " Longhorn " , given to developers at the Microsoft Professional Developers Conference in 2003 , included WinFS , but it suffered from significant performance issues . In August 2004 , Microsoft announced that WinFS would not ship with Windows Vista ; it would instead be available as a downloadable update after Vista 's release .

On August 29 , 2005 , Microsoft quietly made Beta 1 of WinFS available to MSDN subscribers . It worked on Windows XP , and required the .NET Framework to run . The WinFS API was included in the System.Storage namespace . The beta was refreshed on December 1 , 2005 to be compatible with version 2 @.@ 0 of the .NET Framework . WinFS Beta 2 was planned for some time later in 2006 , and was supposed to include integration with Windows Desktop Search , so that search results include results from both regular files and WinFS stores , as well as allow access of WinFS data using ADO.NET.

However , on June 23 , 2006 , the WinFS team at Microsoft announced that WinFS would no longer be delivered as a separate product , and some components would be brought under the umbrella of other technologies - like the object @-@ relational mapping components into ADO.NET Entity Framework ; support for unstructured data , adminless mode of operation , support for file system objects via the FILESTREAM data type , and hierarchical data in SQL Server 2008 , then codenamed Katmai , as well as integration with Win32 APIs and Windows Shell and support for traversal of hierarchies by traversing relationships into later releases of Microsoft SQL Server ; and the synchronization components into Microsoft Sync Framework . However , having a shared @-@ schema storage system built into a future iteration of Microsoft Windows has not yet been ruled out .

With that announcement , most analysts assumed that the WinFS project was being killed off . But in November 2006 , Steve Ballmer said in an interview that WinFS is being actively developed but integration into the Windows codebase will come only after the technology has fully incubated . It was subsequently confirmed in an interview with Bill Gates and that Microsoft plans to migrate applications like Windows Media Player , Windows Photo Gallery , Microsoft Office Outlook etc. to use WinFS as the data storage back @-@ end .

In 2013 Bill Gates cited WinFS as his greatest disappointment at Microsoft and that the idea of WinFS was ahead of its time , which will re @-@ emerge .

= = Data storage = =

= = = Architecture = = =

WinFS uses a relational engine , which is derived from SQL Server 2005 , to provide the data relations mechanism . WinFS stores are simply SQL Server database ( .MDF ) files with the FILESTREAM attribute set . These files are stored in access @-@ restricted folder named " System Volume Information " placed into the volume root , in folders under the folder " WinFS " with names of GUIDs of these stores .

At the bottom of the WinFS stack lies WinFS Core which interacts with the filesystem and provides file access and addressing capabilities . The relational engine leverages the WinFS core services to present a structured store and other services such as locking which the WinFS runtime uses to implement the functionality . The WinFS runtime expose Services such as Synchronization and Rules which can be used to synchronize WinFS stores or perform certain actions on the occurrence of certain events .

WinFS runs as a service which runs three processes - WinFS.exe , which hosts relational datastore , WinFSSearch.exe , which hosts the indexing and querying engine , and WinFPM.exe ( WinFS File Promotion Manager ) , which interfaces with the underlying file system . It allows programmatic access to its features , via a set of .NET Framework APIs , that enables applications to define custom made data types , define relationships among data , store and retrieve information , and allow advanced searches . The applications can then aggregate the data and present the aggregated data to the user .

= = = Data store = = =

WinFS stores data in relational stores , which are exposed as virtual locations called stores . A WinFS store is a common repository where any application can store data along with its metadata , relationships and schema . WinFS runtime can apply certain relationships itself ; for example , if the values of the subject property of a picture and the name property of a contact are same , then WinFS can relate the contact with the picture . Relations can also be specified by other applications or the user .

WinFS provides a unified storage , but stops short of defining the format that is to be stored in the data stores . Instead it supports data to be written in application specific formats . But applications must provide a schema that defines how the file format should be interpreted . For example , a schema could be added to allow WinFS to understand how to read and thus be able to search and analyze , say , a PDF file . By using the schema , any application can read data from any other application , and also allows different applications to write in each other ' s format by sharing the schema .

Multiple WinFS stores can be created on a single machine . This allows different classes of data to be kept segregated ; for example , official documents and personal documents can be kept in different stores . WinFS , by default , provides only one store , named " DefaultStore " . WinFS stores are exposed as shell objects , akin to Virtual folders , which dynamically generates a list of all items present in the store and presents them in a folder view . The shell object also allows searching information in the datastore .

A data unit that has to be stored in a WinFS store is called a WinFS Item . A WinFS item , along with the core data item , also contains information on how the data item is related to other data . This Relationship is stored in terms of logical links . Links specify which other data items the current item is related with . Put in other words , links specify the relationship of the data with other data items . Links are physically stored using a link identifier , which specifies the name and intent of the relationship , such as type of or consists of . The link identifier is stored as an attribute of the data item . All the objects which have the same link id are considered to be related . An XML schema ,

defining the structure of the data items that will be stored in WinFS , must be supplied to the WinFS runtime beforehand . In Beta 1 of WinFS , the schema assembly had to be added to the GAC before it could be used .

== Data model ==

WinFS models data using the data items , along with their relationships , extensions and rules governing its usage . WinFS needs to understand the type and structure of the data items , so that the information stored in the data item can be made available to any application that requests it . This is done by the use of schemas . For every type of data item that is to be stored in WinFS , a corresponding schema needs to be provided to define the type , structure and associations of the data . These schemas are defined using XML .

Predefined WinFS schemas include schemas for documents , e @-@ mail , appointments , tasks , media , audio , video , and also includes system schemas that include configuration , programs , and other system @-@ related data . Custom schemas can be defined on a per @-@ application basis , in situations where an application wants to store its data in WinFS , but not share the structure of that data with other applications , or they can be made available across the system .

== Type system ==

The most important difference between a file system and WinFS is that WinFS knows the type of each data item that it stores . And the type specifies the properties of the data item . The WinFS type system is closely associated with the .NET framework 's concept of classes and inheritance . A new type can be created by extending and nesting any predefined types .

WinFS provides four predefined base types : Items , Relationships , ScalarTypes and NestedTypes . An Item is the fundamental data object which can be stored , and a Relationship is the relation or link between two data items . Since all WinFS items must have a type , the type of item stored defines its properties . The properties of an Item may be a ScalarType , which defines the smallest unit of information a property can have , or a NestedType , which is a collection of more than one ScalarTypes and / or NestedTypes . All WinFS types are made available as .NET CLR classes .

Any object represented as a data unit , such as contact , image , video , document etc . , can be stored in a WinFS store as a specialization of the Item type . By default , WinFS provides Item types for Files , Contact , Documents , Pictures , Audio , Video , Calendar , and Messages . The File Item can store any generic data , which is stored in file systems as files . But unless an advanced schema is provided for the file , by defining it to be a specialized Item , WinFS will not be able to access its data . Such a file Item can only support being related to other Items .

A developer can extend any of these types , or the base type Item , to provide a type for his custom data . The data contained in an Item is defined in terms of properties , or fields which hold the actual data . For example , an Item Contact may have a field Name which is a ScalarType , and one field Address , a NestedType , which is further composed of two ScalarTypes . To define this type , the base class Item is extended and the necessary fields are added to the class . A NestedType field can be defined as another class which contains the two ScalarType fields . Once the type is defined , a schema has to be defined , which denotes the primitive type of each field , for example , the Name field is a String , the Address field is a custom defined Address class , both the fields of which are Strings . Other primitive types that WinFS supports are Integer , Byte , Decimal , Float , Double , Boolean and DateTime , among others . The schema will also define which fields are mandatory and which are optional . The Contact Item defined in this way will be used to store information regarding the Contact , by populating the properties field and storing it . Only those fields marked as mandatory needs to be filled up during initial save . Other fields may be populated later by the user , or not populated at all . If more properties fields , such as last conversed date , need to be added , this type can be extended to accommodate them . Item types for other data can be defined similarly .

WinFS creates tables for all defined Items . All the fields defined for the Item form the columns of

the table and all instances of the Item are stored as rows in the table for the respective Items . Whenever some field in the table refers to data in some other table , it is considered a relationship . The schema of the relationship specifies which tables are involved and what the kind and name of the relationship is . The WinFS runtime manages the relationship schemas . All Items are exposed as .NET CLR objects , with a uniform interface providing access to the data stored in the fields . Thus any application can retrieve object of any Item type and can use the data in the object , without being aware of the physical structure the data was stored in .

WinFS types are exposed as .NET classes , which can be instantiated as .NET objects . Data are stored in these type instances by setting their properties . Once done , they are persisted into the WinFS store . A WinFS store is accessed using an ItemContext class ( see Data retrieval section for details ) . ItemContext allows transactional access to the WinFS store ; i.e. all the operations since binding an ItemContext object to a store till it is closed either all succeed or are all rolled back . As changes are made to the data , they are not written to the disc ; rather they are written to an in @-@ memory log . Only when the connection is closed are the changes written to the disc in a batch . This helps to optimize disc I / O. The following code snippet , written in C # , creates a contact and stores it in a WinFS store .

===== Relationships =====

A datum can be related to one more item , giving rise to a one @-@ to @-@ one relationship , or with more than one items , resulting in a one @-@ to @-@ many relationship . The related items , in turn , may be related to other data items as well , resulting in a network of relationships , which is called a many @-@ to @-@ many relationship . Creating a relationship between two Items create another field in the data of the Items concerned which refer the row in the other Item ? s table where the related object is stored .

In WinFS , a Relationship is an instance of the base type Relationship , which is extended to signify a specialization of a relation . A Relationship is a mapping between two items , a Source and a Target . The source has an Outgoing Relationship , whereas the target gets an Incoming Relationship . WinFS provides three types of primitive relationships ? Holding Relationship , Reference Relationship and Embedding Relationship . Any custom relationship between two data types are instances of these relationship types .

Holding Relationships specifies ownership and lifetime ( which defines how long the relationship is valid ) of the Target Item . For example , the Relationship between a folder and a file , and between an Employee and his Salary record , is a Holding Relationship ? the latter is to be removed when the former is removed . A Target Item can be a part of more than one Holding Relationships . In such a case , it is to be removed when all the Source Items are removed .

Reference Relationships provide linkage between two Items , but do not have any lifetime associated , i.e. , each Item will continue to be stored even without the other .

Embedding Relationships give order to the two Items which are linked by the Relationship , such as the Relationship between a Parent Item and a Child Item .

Relationships between two Items can either be set programmatically by the application creating the data , or the user can use the WinFS Item Browser to manually relate the Items . A WinFS item browser can also graphically display the items and how they are related , to enable the user to know how their data are organized .

===== Rules =====

WinFS includes Rules , which are executed when a certain condition is met . WinFS rules work on data and data relationships . For example , a rule can be created which states that whenever an Item is created which contains field " Name " and if the value of that field is some particular name , a relationship should be created which relates the Item with some other Item . WinFS rules can also access any external application . For example , a rule can be built which launches a Notify application whenever a mail is received from a particular contact . WinFS rules can also be used to

add new properties fields to existing data Items .

WinFS rules are also exposed as .NET CLR objects . As such any rule can be used for any purpose . A rule can even be extended by inheriting from it to form a new rule which consists of the condition and action of the parent rule plus something more .

== RAV ==

WinFS supports creating Rich Application Views ( RAV ) by aggregating different data in a virtual table format . Unlike database view , where each individual element can only be a scalar value , RAVs can have complex Items or even collections of Items . The actual data can be across multiple data types or instances and can even be retrieved by traversing relationships . RAVs are intrinsically paged ( dividing the entire set of data into smaller pages containing disconnected subsets of the data ) by the WinFS runtime . The page size is defined during creation of the view and the WinFS API exposes methods to iterate over the pages . RAVs also supports modification of the view according to different grouping parameters . Views can also be queried against .

== Access control ==

Even though all data are shared , everything is not equally accessible . WinFS uses the Windows authentication system to provide two data protection mechanisms . First , there is share @-@ level security that controls access to your WinFS share . Second , there is item level security that supports NT compatible security descriptors . The process accessing the item must have enough privileges to access it . Also in Vista there is the concept of " integrity level " for an application . Higher integrity data cannot be accessed by a lower integrity process .

== Data retrieval ==

The primary mode of data retrieval from a WinFS store is querying the WinFS store according to some criteria , which returns an enumerable set of items matching the criteria . The criteria for the query is specified using the OPath query language . The returned data are made available as instances of the type schemas , conforming to the .NET object model . The data in them can be accessed by accessing the properties of individual objects .

Relations are also exposed as properties . Each WinFS Item has two properties , named IncomingRelationships and OutgoingRelationships , which provides access to the set of relationship instances the item participates in . The other item which participates in one relationship instance can be reached through the proper relationship instance .

The fact that the data can be accessed using its description , rather than location , can be used to provide end @-@ user organizational capabilities without limiting to the hierarchical organization as used in file @-@ systems . In a file system , each file or folder is contained in only one folder . But WinFS Items can participate in any number of holding relationships , that too with any other items . As such , end users are not limited to only file / folder organization . Rather , a contact can become a container for documents ; a picture a container for contacts and so on . For legacy compatibility , WinFS includes a pseudo @-@ type called Folder which is present only to participate in holding relationships and emulate file / folder organization . Since any WinFS Item can be related with more than one Folder item , from an end user perspective , an item can reside in multiple folders without duplicating the actual data . Applications can also analyze the relationship graphs to present various filters . For example , an email application can analyze the related contacts and the relationships of the contacts with restaurant bills and dynamically generate filters like " Emails sent to people I had lunch with " .

== Searches ==

The WinFS API provides a class called the ItemContext class , which is bound to a WinFS store .

The ItemContext object can be used to scope the search to the entire store or a subset of it . It also provides transactional access to the store . An object of this class can then spawn an ItemSearcher object which then takes the type ( an object representing the type ) of the item to be retrieved or the relationship and the OPath query string representing the criteria for the search . A set of all matches is returned , which can then be bound to a UI widget for displaying en masse or enumerating individually . The properties items can also be modified and then stored back to the data store to update the data . The ItemContext object is closed ( which marks the end of association of the object with the store ) when the queries are made or changes merged into the store .

Related items can also be accessed through the items . The IncomingRelationships and OutgoingRelationships properties give access to all the set of relationship instances , typed to the name of the relationship . These relationship objects expose the other item via a property . So , for example , if a picture is related to a picture , it can be accessed by traversing the relationship as :